Memory error card

Controlled speaker filter · gpt-oss-120b · random links at d = 0.8
200 workplace questions / 5 histories; 80 social questions / 41 histories.
Three draws and four option orders. Message text and names stay fixed.
Comparator: correct index. Budget: 2,048 tokens; shortlist: 64 messages.

Measure	Workplace	Social
Exposure	95.50%	84.08%
Context change	81.00%	73.37%
Evidence displacement	22.83%	3.25%
Churn	19.25% churn. Gains 6.04%, losses 7.79%. Net -1.75 points.	4.07% churn. Gains 1.05%, losses 2.57%. Net -1.52 points.
Conditional change	Loss among right: 14.43%. Gain among wrong: 13.13%.	Loss among right: 3.20%. Gain among wrong: 5.38%.
Index headroom	-0.62 points	+0.51 points
No-memory agreement	63.90% of stable valid answers after context change. Stratum mass 61.75%.	29.74% of stable valid answers after context change. Stratum mass 69.31%.
Repeated requests	3/24 selected repeated requests disagree.	5/24 selected repeated requests disagree.
Loss concentration	75/200 questions have a loss. Top 20: 47.59% of loss; top question: 4.28%.	5/80 questions have a loss. Top 8: 100.00% of loss; top question: 78.95%.

Rates retain equal history weights and original mass before conditioning.
Churn includes service variation and one invalid workplace answer.
Unchanged contexts reuse answers. Net = gain minus loss.
Headroom does not bound harm. No-memory agreement does not prove non-use.
Replay counts use selected requests across conditions, not this full population.
Synthetic histories; imperfect evidence labels; descriptive saved-data analyses.
Point estimates only. The paper gives social history-bootstrap intervals and workplace omit-one-history sensitivity ranges.
Only 5 social questions lose answers, so the top-8 share must be 100%.
Concentration uses weighted losses; its independent simulation omits dependence.